

# Team 2 – Usagi

---

# Team Members

Organization of all team members: National Taiwan University



Tso-Fei Yen



Jui-Chien Tsou



Wei-Chin Wang



Kuan-Hsun Tu



Hsuan-Chi Liu



Chia-Yi Chin



Hsin-Lu Yeh



Fan-Shi Liu



Jui-En Lee




Prof. Chun-Yi Lee

# Mentors

---



Reese Wang — NVidia



Johnson Sun — NVidia


# Exascale Climate Emulator

---


- Abdulah et al. (2024) introduced the **Exascale Climate Emulator (ECE)**, leveraging **SHT**(Spherical Harmonic Transform) and **Cholesky factorization** to boost emulating resolution and throughput.
- We build on this foundation to optimize and scale the climate emulator on our GPU cluster to improve its performance

# Exascale Climate Emulator–Framework

## Climate Emulator: Development







## Climate Emulator: Emulation



Source: Abdulah, S., et al




# Optimization Strategy

## Stage 1



MPI distribution

## Stage 2



Math Library

Mix Precision Optimization

Pipeline Optimization

Problem size Tuning

Cuda Stream Tuning

**Result: Achieve 777.7x Speedup and 119.8x Energy Efficient**

# Stage 1: Data preprocessing

# Strategy – MPI in Data Pre-processing

---

## Steps


- Read Forcing Data
- Run NetCDF Files
- Run Mean Trend (Longest time!)

## Method

- Add MPI, process each location independently → communication overhead
- CHAMELEON → BLAS/LAPACKE, 7x speedup!

# Result – MPI in Data Pre-processing

---



# Stage 2: Cholesky Decomposition

# CPU Baseline

---

We first test out application a CPU node using the following configuration:


**Setup:** 112 CPU cores. 28 MPI ranks  $\times$  4 threads

**Goal:** establish a reference point before GPU/MP optimizations.

**Result: 3264 GFLOPS**

# Strategy 1 – Cublas to CublasLt


- Nsight Systems shows frequent cublasGemmEx calls.
  - Idea: Replaced it with cublasLtMatmul (cuBLASLt) .



# Result – Cublas to CublasLt

---


cublasLt delivered a 4.5~10% improvement.



# Strategy 2 – Mixed Precision

---

Following Abdullah et al. (2024), we introduce better tile-level mixed precision.




DP



SP



HP



Source: Abdullah, S., et al

# Results – Mixed Precision


---

Achieve more than **2x**  
improvement

## L1 loss accuracy check



- DP/SP/HP :  $5.3e-08$
- DP/HP :  $9.6e-06$

→ Still remain accurate



# Strategy 3 – Data Conversion Overhead


Use cublasLt's internal FP conversions instead of explicit float2half



# Result – Data Conversion Overhead

---

Yielded 11~14% improvement.



# Strategy 4 – Adjust N/NB Size

---

## N: matrix order

- Larger **N** increases work, keeps GPUs busier, and improves throughput.
- But too large **N** may overwhelm CPU RAM

## NB: tile size


- Larger **NB** boost higher arithmetic intensity
- But too large **NB** may increase DP percentage and harm performance

Need to strike a balance to maximize occupancy and intensity

# Results – Adjust N/NB Size

---


NB = 4096 with N = 409600  
performs the best and  
reached **2450 TFLOPS!**



# Strategy 5 – Cuda Streams number tuning

---

- **1 stream:** serialized; low overlap/SM util.




- **4+ streams:** concurrent & higher overlap but may contend for SM resources.



# Results – Cuda Streams number tuning

---

After tuning, **3 cuda streams** strike the best balance between concurrency and resource contention. Achieving **2.47 PFLOPs** throughput.



Experiments are tested under problem size:  $N/NB=409600/4096$

# Final Performance

---

- Implement cuBLASLt
- Mix precision optimization
- Pipeline simplify
- Problem size tuning
- Multi-stream adjustment

Combined above optimizations, we achieve 2.47 PFLOPs, 777x speedup

# Energy Efficiency

INPUTS			
# CPU Cores	112		
# GPUs (H100)	8		
Application Speedup	777.0x		
<b>Node Replacement</b>			
777.0x			
GPU NODE POWER SAVINGS			
	Intel Platinum 8480+	8x H100 80GB	Power Savings
Compute Power (W)	784,770	6,760	778,010
Networking Power (W)	36,081	93	35,988
<b>Total Power (W)</b>	<b>820,851</b>	<b>6,853</b>	<b>813,998</b>
Node Power efficiency			
119.8x			
ANNUAL ENERGY SAVINGS PER GPU NODE			
	Intel Platinum 8480+	8x H100 80GB	Power Savings
Compute Power (kWh/year)	6,874,585	59,218	6,815,368
Networking Power (kWh/year)	316,073	814	315,259
<b>Total Power (kWh/year)</b>	<b>7,190,658</b>	<b>60,031</b>	<b>7,130,627</b>
\$/kWh	0.18		
Annual Cost Savings	\$ 1,283,512.78		
3-year Cost Savings	\$ 3,850,538.34		
Metric Tons of CO2	5,056		
Gasoline Cars Driven for 1 year	1,091		
Seedlings Trees grown for 10 years	83,571		
(source: <a href="#">Link</a> )			

POWER ASSUMPTIONS		
Node Configurations	Baseline Node	Alternative
	Intel Platinum 8480+	8x H100 80GB
CPU SKU	Intel Platinum 8480+	Intel Platinum 8480+
# CPU	2	2
# CPU Cores	112	112
<b>CPU Power (W)</b>	<b>650</b>	<b>650</b>
GPU SKU	0	H100 80GB SXM4
# GPU	0	8
<b>GPU Power (W)</b>	<b>0</b>	<b>5600</b>
Network Type	IB EDR	IB EDR
# Network Ports	2	2
Network Card Power (W)	60	60
RBM Power (W)	300	450
<b>Total Compute Node Power (W)</b>	<b>1010</b>	<b>6760</b>
Core Network Power / Node	46	93
<b>Total Power / Node</b>	<b>1056</b>	<b>6853</b>


## ASSUMPTIONS

- (1) The workload being input will run 24/7/365 on the node in question
- (2) When the workload runs on a fraction of a CPU or GPU server, no other bottlenecks occur to stop it from scaling up to occupy the full server
- (3) The calculations use TDP for both CPU and GPU. In reality, neither server will run full time at TDP. The comparison here is "worst case CPU" vs. "worst case GPU"
- (4) Annual cost savings are operational for electricity only. Capital, personnel, etc are not included
- (5) Perfect scaling of the workload to multiple nodes for CPUs
- (6) Fractional workload scaling for both CPU and GPU nodes
- (7) The GPU machine runs the CPUs at full speed, full power draw

# Energy Efficiency

---

119x




# Future Works

---

Workload imbalance between processes

- CUDA Graphs, matrix fusion



# Problems Encountered During Optimization

---

## Cublas operation does not append to different cuda streams

- Solution: Use `cublasGetStream()` to bind cublas on cuda streams

## Problem size is bottlenecked by host RAM.

- We tune **N** and **NB** to pack GPUs with tiles, without exceeding **CPU RAM**.

## Unstable execution time and throughput

- CPU bind to reduce resource contention.

# Wishlist

---

## What do you wish existed to make your life easier?

48 hrs in a day, and more cores to run the process faster!

## Event

More opportunities to get to know other groups better

## Systems

Each member has their own account instead of a team account

# Final Thoughts

---

## Was this Open Hackathon worth it?

Absolutely, pushed us to make progress every week with tight deadlines and meetings.

## Future plans

Connect every stage, try running the full pipeline, and identify bottlenecks for further improvements.

## What resources/support will be critical for your work after the event?

- H100 machine, back to V100 now :(
- Mentor support

## Application Background

Exascale Climate Emulator addresses the escalating computational and storage requirements of traditional Earth System Model simulations:


- Ultra-high spatial resolution of approximately 3.5 km in space
- Very low memory/computation cost.

## Hackathon Objectives and Approach

Main Objective:

Accelerate Cholesky Factorization

- Cublas and CublasLt
- Mixed-precision computation
- Data conversion overhead
- Different N/NB sizes
- Cuda streams number tuning



**Fig.1 Profiling and Experiment Results.** This application is CUDA-API-heavy, but can be improved by sequence of optimization.

## Technical Accomplishments and Impact

Speedup: 777x!

- Reduces months of computation to just days
- Saves energy equivalent to the emissions of 1,000 cars/year or the absorption of 80,000 trees over a decade.
- Faster climate prediction. Marks a major milestone in climate science and high performance computing.


# Appendix

---

# A storyline for publication on NCHC's website.

高解析度大氣模擬計算器


研究領域示意圖



台大 Usagi 團隊來自李濬屹老師帶領的 ElsaLab 實驗室，將大氣模擬計算器加速了 777 倍！！

Exascale climate emulator 是新一代氣候模擬技術，它讓我們能以接近真實地球規模的速度與精度模擬氣候變化。傳統的高解析度氣候模型需要龐大的計算資源與時間，常常一個模擬要跑上數月甚至數年。而我們的系統透過 GPU 混合精度運算與分散式架構，大幅縮短模擬時間，達到 exascale(百億億次運算)級效能。這樣的加速代表著原本需要幾個月的運算現在可能只需幾天，節省的能量相當於 5,056 公噸二氧化碳，也就是一千多輛汽車一年的排放或八萬多棵樹十年的吸碳量。這樣的突破讓研究者能更快預測極端氣候、模擬不同政策下的地球變化，並推動永續與減碳決策，是氣候科學與高效能運算領域的重要里程碑。

實驗結果



報告投影片連結(由國網上傳到github)