

2025 NCHC Open Hackthon

Optimizing Motion Capture and Athlete Reconstruction
Systems through Multi-View 3D Reconstruction

Team 9 Paw Patrol

Paw Patrol

- Team Members from NTHU IGVI Lab



Tsung-Hsun Tsai



Calvin Ku



Pu Ching



Shang-Ching Liu



Prof. Hung-Kuo Chu



Johnson Sun



Zih-Syuan Jhan



Chen-Hao Huang



Tito Aribowo



Prof. Min-Chun Hu

- Mentor from Nvidia

Optimizing Motion Capture and Athlete Reconstruction Systems through Multi-View 3D Reconstruction

Introduction and Motivation:

Human motion capture and 3D reconstruction are key technologies for understanding and analyzing human movements in sports. Reconstructing accurate **human meshes** from multi-view images provides detailed 3D geometry beyond traditional 2D/3D pose estimation, enabling biomechanical analysis and realistic digital human generation.

However, current reconstruction pipelines are **computationally expensive** and **too slow** for real-time or large-scale applications.



Original Reconstruction Pipeline


The pipeline (based on existing mocap framework—Easymocap) performs multi-view 3D human pose reconstruction using a combination of deep learning for human detection and 2D keypoint estimation, followed by optimization-based human mesh (SMPL) fitting for each detected person.

```
[Timing] Total time per step stage:  
detect : 1233.250s  
keypoints2d : 2629.981s  
vis_2d : 0.137s  
match : 96.607s
```

Keypoint3d Profiling


```
[Timing] Total time per final stage:  
write_raw : 7.738s  
collect : 0.096s  
load_body_model : 0.000s  
fitting_each_person : 4338.863s  
write : 111.342s
```

fitSMPL Profiling




Sequential estimation!

Problem to Solve



per-frame sequential process !



per-subject sequential process !


Goal and Strategy

- What was your goal for coming here?

Accelerate the multi-view 3D human-pose reconstruction pipeline without affecting final reconstruction quality.

- What was your initial strategy?


1. Batched process for detect and keypoint2d
2. Parallel processing the fitsmpl stage
3. Optimization of the SMPL fitting stage



Evaluation

- What were you accomplished?

1. Batched process for Detect and Keypoint2d



Evaluation



- What were you accomplished?
 2. Multiprocessing in SMPL fitting

```
# Each worker preloads SMPL fitting stages
def init_worker(stage_args):
    STAGES = {k: load_model(v) for k, v in stage_args.items()}

# Each worker optimizes one person through all stages
def worker(pid, data):
    for stage in STAGES.values():
        data = stage.optimize(data)
    return pid, data

# Main process: persistent pool for parallel execution
pool = Pool(initializer=init_worker, initargs=(stages,))
results = pool.map(worker, persons)
```

Pseudo Code




Evaluation

fitSMPL:
1.init_params
2.fitShape
3.init_RT
4.refine_poses

- What were you accomplished?


3. Reduced LBFGS history size (100->10)



Time
consuming !

```
q = flat_grad.neg()
for i in range(num_old - 1, -1, -1):
    al[i] = old_stps[i].dot(q) * ro[i]
    # q.add_(-al[i], old_dirs[i])
    q.add_(old_dirs[i], alpha=-al[i])

# multiply by initial Hessian
# r/d is the final direction
d = r = torch.mul(q, H_diag)
for i in range(num_old):
    be_i = old_dirs[i].dot(r) * ro[i]
    # r.add_(al[i] - be_i, old_stps[i])
    r.add_(old_stps[i], alpha=(al[i] - be_i))
```



Experiments

We considered two cases

1. Single-person (simplest case)

- Evaluate improvements in each stage: **Detection**, **Keypoints2D**, **SMPL**
- Serves as the **baseline** for per-stage efficiency.

baseball 3views



2. Multi-person (general case)

- Each stage's workload scales with the number of people, making the original pipeline much slower.
- Tested **per-stage** and **overall** performance gains after optimization.

boxing 3views



Results and Final Profile

Strategies

1. batched detect and keypoint2d
2. Parallel process the fitsmpl stage
3. reduce lbfgs history size

Case1: baseball0 795 frames 3views - Single-person

Overall x8.5, Detect x12.5, Keypoint2d x20.8, fitsmpl x4.6 in one person case (compare with Baseline on RTX 2060)

Overall x2.2, Detect x3.7, Keypoint2d x5.3, fitsmpl x1.5 in one person case (compare with Baseline on A100)

Settings (test on single-person data)	CPU device	device	Detect	Keypoint2d	Fitsmpl	total time
Baseline	Intel(R) Core(TM) i7-9700K CPU @ 3.60GHz	RTX 2060	143.2 s	239.3 s	94.0 s	655.0 s
Baseline	AMD EPYC 7742 64-Core Processor	A100	42.9 s (3.3x)	61.0 s (3.9x)	30.1 s (3.1x)	167.3 s (3.9x)
Stratagy1 (batch size = 3*32)	AMD EPYC 7742 64-Core Processor	A100	11.5 s (12.5x)	11.5 s (20.8x)	30.1 s (3.1x)	86.1 s (7.6x)
Stratagy1 (batch size = 3*32) + Stratagy3	AMD EPYC 7742 64-Core Processor	A100	11.5 s (12.5x)	11.5 s (20.8x)	20.5 s (4.6x)	77.2 s (8.5x)

Results and Final Profile

Case2: boxing 600 frames 3views - Multi-person

Strategies

1. batched detect and keypoint2d
2. Parallel process the fitsmpl stage
3. reduce lbfgs history size

Overall x6.4, Detect x13.9, Keypoint2d x8.5, fitsmpl x6.2 in multi-person case (compare with Baseline on RTX 2060)

Overall x2.5, Detect x7.6, Keypoint2d x2.8, fitsmpl x2.7 in multi-person case (compare with Baseline on A100)

Settings (test on multi-person data 600 frames)	CPU device	device	Detect	Keypoint2d	Fitsmpl	Total time
Baseline	Intel(R) Core(TM) i7-9700K CPU @ 3.60GHz	RTX 2060	130.2 s	316.4 s	738.5 s	1364.6 s
Baseline	AMD EPYC 7742 64-Core Processor	A100	71.1 s (1.8x)	103.4 s (3.1x)	321.2 s (2.3x)	529.9 s (2.6x)
Stratagy1 (batch size = 3*32) + Stratagy3	AMD EPYC 7742 64-Core Processor	A100	9.4 s (13.9x)	37.4 s (8.5x)	222.8 s (3.3x)	310.2 s (4.4x)
Stratagy1 (batch size = 3*32) + Stratagy2 (Multi-Threading version) + Stratagy3	AMD EPYC 7742 64-Core Processor	A100	9.4 s (13.9x)	37.4 s (8.5x)	146.4 s (5.0x)	228.4 s (6.0x)
Stratagy1 (batch size = 3*32) + Stratagy2 (Multi-Processing version) + Stratagy3	AMD EPYC 7742 64-Core Processor	A100	9.4 s (13.9x)	37.4 s (8.5x)	119.6 s (6.2x)	211.6 s (6.4x)

Final Thoughts

- Was this Open Hackathon worth it?

Yes. **Special thanks to mentor, Johnson Sun**

- What did you learn?

Profiling Cycle


- Future Works

Optimization of data loading

TensorRT (detect + keypoints)

Optimization of LBFGS approximate inverse Hessian process

Use neural network-based initialization to provide better SMPL parameters and reduce optimization iterations.



Create a storyline for publication on NCHC's website.

3D人體姿勢重建

Team 9 : Prof. Hung-Kuo Chu, Prof. Min-Chun Hu, Tsung-Hun Tsai, Calvin Ku, Pu Ching, Shang-Ching Liu, Zih-Syuan Jhan, Chen-Hao Huang, Tito Aribowo

per-frame multiview images (frame x views) → batched → Keypoint3d:
1. detect
2. keypoints2d
3. match
4. triangulate → per-frame 3D keypoints → MultiProcessing → fitSMPL:
1. init_params
2. fitShape
3. init_RT
4. refine_poses → per-frame SMPL

Paw-Patrol 團隊成員來自清華大學資工系「朱宏國和胡敏君老師」帶領的IGVI實驗室，將3D人體姿勢重建加速了 8.5 倍！— NVIDIA Mentor: Johnson Sun

Table1: Speedup in single person case

Settings	device	Detect	Keypoint2d	Fitsmpl	total time
Baseline	RTX 2060	143.2 s	239.3 s	94.0 s	655.0 s
Baseline	A100	42.9 s (3.3x)	61.0 s (3.9x)	30.1 s (3.1x)	167.3 s (3.9x)
Optimized	A100	11.5 s (12.5x)	11.5 s (20.8x)	20.5 s (4.6x)	77.2 s (8.5x)

Table2: Speedup in multi person case

Settings	device	Detect	Keypoint2d	Fitsmpl	Total time
Baseline	RTX 2060	130.2 s	316.4 s	738.5 s	1364.6 s
Baseline	A100	71.1 s (1.8x)	103.4 s (3.1x)	321.2 s (2.3x)	529.9 s (2.6x)
Optimized	A100	9.4 s (13.9x)	37.4 s (8.5x)	119.6 s (6.2x)	211.6 s (6.4x)

Poster Link:

https://www.canva.com/design/DAG4Ym4iswc/5tBFc2CT-FuCZ6W1ZERdHQ/edit?utm_content=DAG4Ym4iswc&utm_campaign=designshare&utm_medium=link2&utm_source=sharebutton

OpenACC
More Science, Less Programming

OPEN
HACKATHONS

NCHC 國家實驗研究院
國家高速網路與計算中心
National Center for High-performance Computing

NVIDIA



Thank You

OpenACC
More Science, Less Programming