



NemoClaw

May 2026

Virginia Chen– Solutions Architect

Today's Agenda

NemoClaw — Run Autonomous Agents Safely

- Part 0 — Market Evolution: from ChatGPT to OpenClaw
- Part 1 — The Problem: Harness & 5 Root Risks
- Part 2 — OpenShell: the Kernel for Agents
- Part 3 — NemoClaw: 3-Layer Stack (Kernel · Distro · Installer)
- Part 4 — Landing: Use Cases & DGX Hardware
- Part 5 — Q&A

The Evolution of AI Agents

From LLMs and multi-turn prompts to long-running agents that modify themselves

ChatGPT

Nov 2022



Foundation LLMs,
Multi-turn Prompts

DeepSeek

Jan 2025




Reasoning Enables Reflection
For Better Outcomes

100x

OpenClaw


Jan 2026



Self Evolving Agents
Working on Your Behalf

1000x

The ChatGPT Moment of Agents Has Arrived



ubuntu@brev-ae6ah8vzm: ~

→ ~ curl -fsSL https://nvidia.com/nemoclaim.sh | bash

→ ~ nemoclaim onboard

The OpenClaw Moment: Unbounded Autonomy



OpenClaw

THE AI THAT ACTUALLY DOES THINGS.


Clears your inbox, sends emails, manages your calendar, checks you in for flights.

All from WhatsApp, Telegram, or any chat app you already use.

NEW

Introducing OpenClaw →

Autonomous Agents Are Multi-Agent AI Systems that Create and Manage Their Own Sub-Agents



Focus on Agent Harnesses



deepagents

To be Useful, Agents Need Autonomy. To Be Safe, They Need Zero Trust. Current Systems Come Up Short.

Sensitive Data Processed

Live API keys held inside the agent. Any compromise leaks data.

Prompt Injection Unbounded

Malicious input takes over – and the agent has full access to act on it.

Self-Evolving Behavior

What the agent can do tomorrow is not what you authorized today.


No Governance for Long Runs

Agents running for hours no oversight, no audit trail, no approvals.


Policy Gaps across Layers

Each layer looks secure in isolation. The interaction between layers is not verified.

Five Core Primitives. Everything Else Is Pluggable.




ACCELERATED TOOLS



NVIDIA OpenShell - The Safer Runtime for Agents

Safely Build and Run Autonomous, Self Evolving Agents



Announcing NVIDIA NemoClaw for OpenClaw



OpenClaw

Definition

Runtime

Memory

Skills

Terminal



NemoClaw


OpenShell



Network Guardrail




Policy Engine



Privacy Router



Local Frontier Model
Nemotron 3 Super



NVIDIA NemoClaw

Open Reference Blueprints for Enterprise Grade Agents



Agent Harness




OpenShell




Open Models




 Policy Engine

 Network Guardrail

 Privacy Router

 Sandbox Supervisor



Nemotron 3



Built in NeMo Acceleration and Optimization

NVIDIA NemoClaw Workflow

Getting started is easy




```
$ curl -fsSL https://www.nvidia.com/nemoclave.sh | bash
```

```
version: 1 # Schema version. Always 1 today.

# ----- STATIC section (locked at sandbox creation) -----
filesystem_policy:
  include_workdir: true # Auto-grants /sandbox as read-write
  read_only: [/usr, /lib, /proc, /dev/urandom, /app, /etc, /var/log]
  read_write: [/sandbox, /tmp, /dev/null]
landlock:
  compatibility: best_effort # or hard_requirement - aborts on failure
process:
  run_as_user: sandbox # root / uid 0 is rejected outright
  run_as_group: sandbox

# ----- DYNAMIC section (hot-reloadable via 'openshell policy update') -----
network_policies:
  github_api:
    name: github-api-readonly
    endpoints:
      - host: api.github.com
        port: 443
        protocol: rest # Triggers L7 HTTP inspection
        enforcement: enforce # or audit - logs but allows
        access: read-only # → only GET, HEAD, OPTIONS
  binaries:
    - { path: /usr/bin/curl } # Only this binary is allowed to use this rule
```



NVIDIA NemoClaw

Blueprint + Installer

Guided onboarding

Step by step wizard installs Node, OpenShell, NemoClaw CLI, creates a sandbox, wires inference, applies policies and drops you into the openclaw TUI

Hardened, digest verified blueprint

Every install is tamper-evident and reproducible

Strict-by-default network policy

Start with a hardened default with scoped network policies to existing endpoint groups

Prewired to Nemotron

Ships with Nemotron 3 Super with the option to switch models as runtime, no restart

Regulated Channel Bridges

Pre-scaffolds outbound channel bridges including Telegram, Discord, and Slack

Why Build NemoClaw?

Enterprise Packaging for Autonomous Agents

“NemoClaw is an opinionated implementation of OpenClaw for the enterprise. No enterprise should deploy OpenClaw. Every enterprise should deploy NemoClaw.”

-- Internal EStaff Feedback / Mission framing in response to JHH

Why Now?

Migration from application specific AI towards **localized agents as the new UX** of the PC



Why NemoClaw?

Agent Harnesses alone is not enough for enterprise deployment. You need **runtime governance, model recommendations, and installation management.**



Why OpenShell?

OpenShell is the **trust boundary** for deployment of long running, autonomous agents.




NemoClaw Use Cases

How could enterprises use NemoClaw?



Personal Assistant/CRM



Secure Coding
Infrastructure



Support + Operations




Personalized Onboarding
Experience




Community
Sentiment/Issue Tracker

Get Started with NVIDIA NemoClaw



github.com/NVIDIA/NemoClaw



docs.nvidia.com/nemoclave/latest

Mini-Demo: 只能看 Public 帳本的記帳助手


Showcasing NemoClaw sandbox host-isolation

- 場景:兩份 CSV — public(日常消費) + private(薪資、投資、匯款)
- 目標:讓 Agent 幫你分類上月花費,但物理上看不到 private
- 機制:NemoClaw sandbox 預設與 host 完全隔離,private 從未進 sandbox = 不可能洩漏
- 差異:這不是 prompt engineering,是 OS-level isolation (Landlock + Seccomp)
- 同樣 pattern 適用:醫療紀錄、合約、產品 roadmap、源碼 — 你決定哪些檔案進 sandbox

Questions around NemoClaw

How are we deploying this internally?

Deploying a community sentiment agent to coordinate roadmap features for NVIDIA products based on public asks and internal priorities



Enterprise vs Open-Source Support

Both OpenShell and NemoClaw are alpha, but working fast for Computex announcements with partners who would provide enterprise support



Support for Multi Tenancy?

This is a highly requested feature that will be done in coordination with partners, including RedHat and SAP. Additional driver support is being added as we speak.



